

Product Requirements Document (PRD)

Agentible Nexus MVP

Overview

Build the first MVP of **Agentible Nexus**, an AI infrastructure product developed by Agentible.

Agentible Nexus is the AI Control Plane for Engineering Teams.

Instead of every application communicating directly with different LLM providers, every request flows through Nexus. The platform intelligently routes requests to the appropriate model, caches repeated requests, logs usage, estimates costs, and gives engineering leaders visibility into AI consumption.

The goal of this MVP is **not** to build a complete enterprise platform.

The goal is to build a polished, production-quality MVP that demonstrates the core value proposition and captures early customer interest through a waitlist.

This project should be designed as if it were the first product of a long-term AI infrastructure company.

Product Vision

Every engineering organization will eventually have dozens of AI-powered applications.

Without an infrastructure layer, companies struggle with:

- Rising LLM costs
- No visibility into AI spending
- Engineers selecting expensive models unnecessarily
- Duplicate requests
- Lack of governance
- Vendor lock-in

Agentible Nexus becomes the intelligent layer between applications and AI providers.

Long-term vision:

Become the operating layer for enterprise AI infrastructure.

Target Users

Primary:

- Engineering Managers
- VP Engineering
- CTOs
- AI Engineering Teams
- Platform Engineers

Secondary:

- Startup founders
 - AI consultants
 - Companies building internal AI tools
-

Core MVP Goals

The MVP must demonstrate these five capabilities:

1. Intelligent AI Gateway
2. Model Routing
3. Request Caching
4. Spend Visibility
5. Waitlist Collection

Everything else is out of scope.

Tech Stack

Framework

- Next.js 15
- TypeScript
- App Router

UI

- Tailwind CSS
- shadcn/ui
- Lucide Icons

Backend

- Next.js API Routes

Database

- Supabase

Caching

- Upstash Redis

AI

- OpenAI SDK

Deployment

- Vercel

Authentication

None for MVP.

Pages

Home (/)

The homepage should immediately communicate:

"What is Agentible Nexus?"

Hero Section

Headline:

The AI Control Plane for Engineering Teams

Subheading:

Route every AI request through one intelligent gateway. Optimize model selection, reduce inference costs, improve cache efficiency, and gain complete visibility into AI usage across your engineering organization.

Primary CTA

Join the Waitlist

Secondary CTA

View Interactive Demo

Sections

Problem

Explain why AI infrastructure is becoming essential.

Pain points:

- AI spending grows unpredictably.
- Every team chooses different models.
- Expensive requests go unnoticed.
- Duplicate prompts waste money.
- Organizations lack visibility.

Solution

Explain how Nexus solves these issues.

Illustrate:

Application

↓

Agentible Nexus

↓

Routing Engine

↓

Cache

↓

Selected LLM

↓

Analytics Dashboard

Features

Display cards for:

- AI Gateway
- Smart Model Routing
- Prompt Caching
- Spend Analytics
- Centralized Governance

Built For

Cards for:

- Startups
- Enterprise Teams
- AI Products
- Internal Engineering Platforms

CTA

Large waitlist section encouraging visitors to join.

About (/about)

Explain the company.

Agentible is an AI Infrastructure company.

Mission:

Build the foundational infrastructure layer that helps organizations deploy, optimize, govern, and scale AI systems.

Explain that Nexus is the company's first product.

Future vision:

- AI Gateway
- AI Analytics
- AI Governance
- AI Security
- AI Evaluation

Demo (/demo)

This page demonstrates the product.

It should contain:

Mock Architecture Diagram

Application

↓

Gateway

↓

Routing

↓

Redis Cache

↓

LLM

↓

Dashboard

Below it:

Interactive Playground.

Users can:

Enter Prompt

Select Task Type

Click Send

Receive Response

Display:

Model selected

Cache hit

Estimated cost

Response time

Dashboard (/dashboard)

Dashboard should feel like an enterprise SaaS.

Widgets

Total Requests

Total Tokens

Estimated Spend

Estimated Savings

Cache Hit Rate

Average Cost per Request

Charts

Requests over time

Spend by Model

Model Usage

Recent Requests Table

Columns

Timestamp

Task

Model

Tokens

Cost

Cache Hit

Use real Supabase data if available.

Otherwise display realistic demo data.

API

POST /api/gateway

Input

Prompt

Task Type

User Email

Flow

Receive request.

Determine task type.

Select appropriate model.

Create cache key.

Check Redis.

If cache exists:

Return cached response.

Log cache hit.

If cache does not exist:

Call OpenAI.

Store response.

Save metadata in Supabase.

Return response.

Task Routing Rules

Planning

Architecture

Reasoning

↓

GPT-4o

Coding

Debugging

Refactoring

↓

GPT-4o-mini

Summaries

Documentation

Extraction

Classification

↓

GPT-4o-mini

Response

Response

Model Used

Estimated Cost

Cache Hit

Execution Time

POST /api/waitlist

Accept

Name

Email

Company

Role

Team Size

Current AI Provider

Biggest AI Cost Challenge

Validate input.

Store in Supabase.

Handle duplicate emails gracefully.

Return success response.

Database

Table

waitlist

Fields

id

name

email

company

role

team_size

provider

challenge

created_at

Table

ai_requests

Fields

id

user_email

task_type

model

prompt

prompt_tokens

completion_tokens

total_tokens

estimated_cost

cache_hit

latency_ms

created_at

Redis

Create cache key from:

model

taskType

prompt

Store:

AI response

Timestamp

TTL

Default TTL

24 hours

UI Components

Navbar

Hero

Feature Cards

Architecture Section

Statistics Cards

Dashboard Widgets

Playground

Waitlist Form

Footer

Buttons

Badges

Tables

Charts

Toast Notifications

Loading States

Empty States

Navigation

Navbar

Home

About

Demo

Dashboard

Waitlist

GitHub

Footer

Product

About

Documentation (Coming Soon)

Privacy

Contact

Built by Agentible

Branding

Company

Agentible

Product

Agentible Nexus

Tagline

The AI Control Plane for Engineering Teams.

Color Palette

Modern dark theme.

Neutral grays.

White typography.

Blue accent.

Minimal gradients.

Design inspiration:

Modern developer infrastructure platforms such as Stripe, Vercel, Linear, and Clerk.

MVP Success Criteria

The MVP is complete when it includes:

- Professional landing page
- Responsive design
- Functional waitlist

- Working AI Gateway
 - Model routing
 - Redis caching
 - Supabase logging
 - Demo playground
 - Analytics dashboard
 - Clean navigation
 - README with setup instructions
 - Deployable on Vercel without code changes
-

Git Commit Plan

Commit 1

chore: initialize Next.js project

Commit 2

chore: configure Tailwind, shadcn, Supabase and Redis

Commit 3

feat: implement AI gateway with intelligent model routing

Commit 4

feat: add Redis caching layer

Commit 5

feat: implement waitlist API and Supabase persistence

Commit 6

feat: build landing page

Commit 7

feat: build interactive AI playground

Commit 8

feat: build analytics dashboard

Commit 9

feat: add about page and navigation

Commit 10

polish: improve UI, responsiveness and animations

Commit 11

docs: add production-ready README

Important Engineering Principles

- Keep the architecture modular and easy to extend.
- Separate business logic from UI components.
- Use reusable React components throughout.
- Follow clean code and TypeScript best practices.
- Prioritize simplicity over feature completeness.
- Optimize for readability, maintainability, and future extensibility.
- Build every component as though this MVP will evolve into a production enterprise platform.